

Concurrency and Parallelism 2018-19

Project — Parallel Patterns (v1.0)

João Lourenço

October 31, 2018

Abstract

This document describes the project assignment for the course of Concurrency and Parallelism.

1 Introduction

Up to last week we have been learning and discussing about Patterns for Parallel Programming. In this project you are asked to collaborate with another colleague and implement a set of Patterns for Parallel Programming using Cilk+. You are supposed to implement at least the following patterns: Map, Reduce, Scan, Pack, Gather, Scatter, Pipeline, Farm, and optimize your code to be as fast and as scalable as possible.

1.1 Working rules

The following working rules apply. This list is not exhaustive, so in case of doubt, do not hesitate in asking.

- The project work will be done in groups of 3 people (unless explicitly authorized otherwise).
- Each group will fork the given main repository and work on their own private repository (remember to keep your repository private *and not public*).
- Each student will commit his/her own work into the shared repository. I expect at least one commit on each day that you work on the project.
- Commit messages must describe clearly what was changed/added in that commit (and why).
- Remember to push your commits to your group repository.
- The project report must have the look and feel of a research paper, using the IEEE template for Computer Society Journals (<https://goo.gl/Xtjdh4>), with a maximum of 5 pages (references may be on the 6th page).

1.2 Project grading aspects

The Project's grade $[\mathcal{P}]$ will take into consideration the following aspects. The order is not relevant and the list is not exhaustive. In case of doubt, do not hesitate in asking.

- The project will be graded according to the following criteria:
 - Projects that fail to compile and execute with the following procedure will receive the final grade of 0 points:

```
clone <your_repository>
cd <your_repository>
make
./main NUMBER
```
 - The quality of the work as perceived from project report;
 - The quality of the text in the project report (well organized, complete, the text/message is clear, no misspellings, etc)
 - The project's code look and feel;
 - A perfect project which implements all (and only) the above listed patterns will have a top grade of 17 (over 20) points.
- Some points are reserved for complementary work/achievements, such as:
 - Some points will be given to the implementation of more patterns (unlisted above). Please discuss with me the *signature* of the patterns you want to implement.
 - Some points will be given to the implementation of more unit/integration functions, if shared publicly with the colleagues and used and acknowledged by them in their reports.
 - Some points will be given to the project that exhibits the best (timed) speedup (compiling and running the s in the *server node9*).
 - Some points will be given to the project that exhibits the best scaled (data size) speedup for 16 processors (compiling and running the s in the *server node9*).

1.3 Student grading aspects

The Students's grade $[\mathcal{S}]$ will take into consideration the following aspects. The order is not relevant and the list is not exhaustive. In case of doubt, do not hesitate in asking.

- Relevance of the student's commit messages.
- Relevance of the code committed by the student.
- Students that do not exhibit evidences of working on the project (e.g., with no commits, or with nearly empty/non-meaningful commits, or with insufficient/unclear commit messages) will fail the course.

1.4 Student/Project Final grading rule

- Formula: $\mathcal{G} = 0.7 * \mathcal{P} + 0.3 * \mathcal{S}$

2 Project Description

In this lab work you are given a working sequential version of a Ray Tracer written using the C++ programming language. You are asked to study the given code and create an optimized (parallel) version of the code using Cilk+.

2.1 Given Version

You are given a (almost working) version of the project at

https://bitbucket.org/cp201819/parallel_patterns.git

This repository contains two directories/folders: `src` and `doc`. The former contains the base source code, and the latter will contain your Project Report as a PDF file. Please name your report as `report_AAAAA_BBBBB_CCCCC.pdf`, where AAAAA, BBBB and CCCCC are the numbers of the group members sorted in increasing order (lowest to highest).

Fork the above repository and name your groups repository as

`cp2018-19_project_AAAAA_BBBBB_CCCCC.pdf`

where AAAAA, BBBB and CCCCC are the numbers of the group members sorted in increasing order (lowest to highest).

In your Linux device (own laptop, lab workstation, or as a last resort, in the “node9” server used in the last lab class) clone your new repository and then try compile your code using the command `make` in the `src` directory. It must compile with no errors nor warnings.

2.2 Code Structure

`debug.c debug.h main.c patterns.c patterns.h unit.c unit.h`

The project include the following source files:

Files	Description
<code>debug.c debug.h</code>	Functions for printing the contents of the array(s), useful for debugging (activated with the option “-d”).
<code>patterns.c patterns.h</code>	The patterns to be implemented. The “.c” file contains empty functions.
<code>unit.c unit.h</code>	Functions for unit testing of each pattern.
<code>main.c</code>	The main program.

2.3 Work plan

Your job is to make an optimized parallel version (using Cilk+) of all the patterns listed in the files `patterns.c/patterns.h`!

You may follow these steps:

1. Clone/fork the given project.
2. Compile the given version. Study the source code and understand how it works.

3. Discuss with your colleagues how to split the work.
4. Implement a sequential version of each pattern.
5. Compile and run the tests and confirm the results.
6. Implement a parallel version of each pattern.
7. Compile and run the tests and confirm the results.
8. For each pattern, measure its performance/scalability/scaled scalability. You may experiment with different numbers of processors (by setting the environment variable `CILK_NWORKERS`).
9. Optimize the given code.
10. Go back to item 7. until satisfied.
11. Write the report. Revise the report. Please put your report in the `doc` directory and name it as `report_AAAAA_BBBBB_CCCCC.pdf`, where AAAAA, BBBBB and CCCCC are the numbers of the group members sorted in increasing order (lowest to highest).
12. *Optional*: implement and optimize some more parallel patterns.
13. *Optional*: implement and share some more tests (unit or integration tests).
14. *Optional*: complete the report and revise again.

Please remember to commit regularly your changes, and please always write meaningful commit messages.

3 Questions/Discussion

Please ask your questions using the Piazza system. Either public (if possible) or private (if really necessary).